

B-Trees I

- definition
- example
- types
- search
- insertion

B-Trees I [Bono]

1

B-Tree introduction

- a search structure, i.e., yet another possible representation for the Table ADT.
- example of a balanced search tree: $O(\log n)$ search time *worst* case.
- other balanced search trees: AVL trees, red-black trees.
- Well suited for *large* amounts of data (on disk), such as for a database application.
(B+ Trees is a variation commonly used)
- Most DBMS use them to represent data.

B-Trees I [Bono]

2

Balanced search structures: main idea

- search tree with extra properties that limit its height (so search will be $\log n$)
- means insert and delete are more complicated, because the resulting trees have to satisfy the extra properties
- but, to keep it fast, insert and delete also need to be $\log n$.


B-Trees I [Bono]

3

B-Tree: definition

A B-Tree of order m is an m -way tree such that:

1. non-leaf nodes store up to $(m-1)$ entries that partition the data in the subtrees in a way similar to a BST (see diagram below)
 - the number of entries is one less than the number of children
 - the entries are sorted from smallest to largest
2. non-leaf nodes (except root): have between $\lceil m/2 \rceil$ and m children.
3. all nodes (except root): have between $\lceil m/2 \rceil - 1$ and $m-1$ entries
4. root (special case):
 - either it's a leaf with at least one entry
 - or it has between 2 and m children.
5. all leaves are at the same depth



B-Trees I [Bono]


4

Example: 5-way B-Tree

- a.k.a. B-Tree of order 5
- $M = 5$
- Main and Savitch call this a B-Tree with $\text{MAXIMUM} = 4$ (i.e., max entries)

nodes (except root)

- 3-5 children (unless leaf)
- 2-4 entries (i.e., $\langle \text{key}, \text{value} \rangle$)



root

- 2-5 children (unless it's a leaf)
- 1-4 keys

B-Trees I [Bono]

5

5-way B-Tree (cont)

Example type definition for a node:

```
const int MAXCHILD = 5;

class Bnode {
private:
    friend class Table;
    int nKeys;
    EntryType entry[MAXCHILD-1];
    Bnode *branch[MAXCHILD];
};


typedef Bnode *BTreeType;
```

B-Trees I [Bono]

6

B-Tree: searching

- keys in interior nodes partition subtrees just like BST nodes do
- Example B-Tree: (keys are letters)



B-Trees I [Bono]

7

B-Tree: insertion

1. search terminates in failure at a leaf.
2. add a new value to leaf in sorted order.
3. if node has more than $M-1$ keys then split the node.

How to split:

1. middle value gets pushed up into the parent
2. first half of values in a new left node.
3. second half of values in right node.

4. If the parent is now bigger than M , split it recursively.

Note₁: if this happens at the root, the root splits in two and get a new root with only 1 key and 2 children. This is the only time we increase the number of levels in the tree.

Note₂: node splits produce two half full nodes; thus we won't have to split again for a while.

B-Trees I [Bono]

8

Insertion: example

- Start with an empty 5-way B-tree and insert the following keys in the order shown (should end up with tree from slide 8):

a g f b k d h m
j e s i r x c l
n t y p

B-Trees I [Bono]

9

B-Tree variations

- 2-3 tree
 - 2-3 tree is a name for the special case of a 3-way B-Tree (all nodes have between 2 and 3 children)
 - Main and Savitch examples where $MAX = 2$ is a 2-3 tree.
- only store data at the leaves
 - another B-Tree formulation is one where interior nodes store keys, but all actual entries are in leaf nodes (more about this later)

B-Trees I [Bono]

10